

## Part 1

DDL Commands:

```
CREATE TABLE Users (username VARCHAR(100) PRIMARY KEY, name  
VARCHAR(100), password VARCHAR(100), is_admin BOOLEAN);
```

```
CREATE TABLE Locations (location_name VARCHAR(250) PRIMARY KEY, address  
VARCHAR(250), city VARCHAR(100), state VARCHAR(100), country VARCHAR(100),  
postal_code VARCHAR(10));
```

```
CREATE TABLE Promoter (promoter_name VARCHAR(100) PRIMARY KEY,  
promoter_url VARCHAR(250));
```

```
CREATE TABLE Events (event_title VARCHAR(250) PRIMARY KEY, event_url  
VARCHAR(250), datetime_local DATETIME, location_name VARCHAR(250),  
promoter_name VARCHAR(100), username VARCHAR(100), FOREIGN KEY  
(location_name) REFERENCES Locations(location_name) ON DELETE CASCADE ON  
UPDATE CASCADE, FOREIGN KEY (promoter_name) REFERENCES  
Promoter(promoter_name) ON DELETE CASCADE ON UPDATE CASCADE,  
FOREIGN KEY (username) REFERENCES Users(username) ON DELETE SET NULL  
ON UPDATE CASCADE);
```

```
CREATE TABLE Tickets (ticket_id VARCHAR(100) PRIMARY KEY, event_title  
VARCHAR(250), ticket_price DECIMAL(10,2), total_price DECIMAL(10,2), fee  
DECIMAL(10,2), full_section VARCHAR(100), section VARCHAR(100), row_num  
VARCHAR(100), quantity INT, username VARCHAR(100), FOREIGN KEY (event_title)  
REFERENCES Events(event_title) ON DELETE CASCADE ON UPDATE CASCADE,  
FOREIGN KEY (username) REFERENCES Users(username) ON DELETE SET NULL  
ON UPDATE CASCADE));
```

```
CREATE TABLE WishList (username VARCHAR(100), event_title VARCHAR(250),  
wishlist_date DATETIME, PRIMARY KEY (username, event_title), FOREIGN KEY  
(username) REFERENCES Users(username) ON DELETE CASCADE ON UPDATE  
CASCADE, FOREIGN KEY (event_title) REFERENCES Events(event_title) ON  
DELETE CASCADE ON UPDATE CASCADE);
```

Users, Tickets, and Wishlist have 1,000+ rows:

```
mysql> SELECT COUNT(*) FROM Users;  
+-----+  
| COUNT(*) |  
+-----+  
|      1000 |  
+-----+  
1 row in set (0.01 sec)
```

```
mysql> SELECT COUNT(*) FROM Tickets;  
+-----+  
| COUNT(*) |  
+-----+  
|    37764 |  
+-----+  
1 row in set (0.00 sec)
```

```
mysql> select count(1) from WishList;  
+-----+  
| count(1) |  
+-----+  
|     1087 |  
+-----+  
1 row in set (0.00 sec)
```

## Advanced SQL Queries:

### **\*\*Events and city with Ticket\_price starting from \$1000\*\***

Select event\_title, city from Events natural join Locations where event\_title in (select event\_title from Tickets where ticket\_price>1000);

```
mysql> Select event_title, city from Events natural join Locations where event_title in (select event_title from Tickets where ticket_price>1000) limit 15;
```

event_title	city
ENHYPEN	Rosemont
Tom Segura	Tampa
Bryan Adams	Orlando
US Open Tennis Championship: Session 3 - Men's/Women's 1st Round	Flushing
Inter Miami CF at Charlotte FC	Charlotte
Hamilton	Fort Myers
Zach Bryan	Brooklyn
SEC Men's Basketball Tournament - Session 6	Nashville
Nashville Predators vs. Colorado Avalanche	Nashville
Ashley McBryde	Durant
Journey & Toto	Billings
O.A.R. & Fitz and The Tantrums	Bridgeport
Netflix Is A Joke Festival - Seinfeld, Gaffigan, Bargatz and Maniscalco	Los Angeles
Alanis Morissette, Joan Jett And The Blackhearts & Morgan Wade	Maryland Heights
Miami Heat vs. Toronto Raptors	Miami

15 rows in set (0.01 sec)

### **\*\*Events and number of tickets available for all events in Great Britain\*\***

select count(ticket\_id), event\_title, event\_url from Tickets natural join Events natural join Locations where country='GB' group by event\_title, location\_name;

```
mysql> select count(ticket_id), event_title, event_url from Tickets natural join Events natural join Locations where country='GB' group by event_title, location_name;
```

count(ticket_id)	event_title	event_url
5	R&B Brunch with Bottomless Rum Punch! (Birmingham)	<a href="https://feverup.com/m/101322">https://feverup.com/m/101322</a>
4	Candlelight: A Tribute to Ludovico Einaudi at Central Hall	<a href="https://feverup.com/m/101185">https://feverup.com/m/101185</a>
15	Candlelight: Hans Zimmer's Best Works at Central Hall Westminster	<a href="https://feverup.com/m/100852">https://feverup.com/m/100852</a>
48	Rum Punch Bottomless Brunch at Cococure	<a href="https://feverup.com/m/100411">https://feverup.com/m/100411</a>
1241	Sherlock: The Official Live Game	<a href="https://feverup.com/m/101498">https://feverup.com/m/101498</a>
28	Top Stand-Up Comedy in Covent Garden	<a href="https://feverup.com/m/100120">https://feverup.com/m/100120</a>
194	Asian Inspired Bottomless Brunch at The Last Talisman	<a href="https://feverup.com/m/100294">https://feverup.com/m/100294</a>

7 rows in set (0.04 sec)

### **\*\*Events, location, city, number of users who have wishlisted events for the first quarter of 2025\*\***

select event\_title, location\_name, city, count(username) from WishList natural join Events natural join Locations where wishlist\_date between '2025-01-01' and '2025-03-01' group by event\_title, location\_name, city;

```
mysql> select event_title, location_name, city, count(username) from WishList natural join Events natural join Locations where wishlist_date between '2025-01-01' and '2025-03-01' group by event_title, location_name, city limit 15;
```

event_title	location_name	city	count(username)
R&B Brunch with Bottomless Rum Punch! (Birmingham)	Bierkeller Birmingham	Birmingham	1
Annie	Washington Pavilion of Arts & Science	Sioux Falls	1
Shane Gillis	The Chicago Theatre	Chicago	1
Willis Alan Ramsey & Keith Sykes	Germantown Performing Arts Centre	Memphis	3
Ryan Beatty	The Wiltern	Los Angeles	1
Augustana	Park West	Chicago	1
Chicago Dogs vs. Cleburne Railroaders	Impact Field	Rosemont	1
Knocked Loose	Marathon Music Works	Nashville	1
Detroit Pistons vs. Toronto Raptors	Little Caesars Arena	Detroit	1
Toughest Monster Truck Tour	Dow Arena At Dow Event Center	Saginaw	2
Shrek The Musical	Louisville Palace	Louisville	2
Oklahoma Sooners vs. BYU Cougars	Lloyd Noble Center	Norman	1
Soul Of Motown	Westgate Cabaret At Westgate Las Vegas Resort & Casino	Las Vegas	2
Big Something	Madison Theater - Covington	Covington	2
Tyler Braden	Tally Ho Theater	Leesburg	1

```
15 rows in set (0.00 sec)
```

## **\*\*Get all Tickets Details, Event details of all events for the top 5 major cities (top 5 cities by events)\*\***

Select ticket\_id, total\_price, event\_title, city from Tickets natural join Events natural join Locations where city in ( Select city from (select city, count(event\_title) from Locations natural join Events group by city order by 2 desc limit 5) z ) limit 15);

```
mysql> Select ticket_id, total_price, event_title, city from Tickets natural join Events natural join Locations where city in ( Select city from (select city, count(event_title) from Locations natural join Events group by city order by 2 desc limit 5) z ) limit 15;
```

ticket_id	total_price	event_title	city
707771229	149.85	Fantasy	Las Vegas
707771232	139.05	Fantasy	Las Vegas
719956804	116.10	Fantasy	Las Vegas
719956883	116.10	Fantasy	Las Vegas
719956898	116.10	Fantasy	Las Vegas
770609924	203.85	Fantasy	Las Vegas
892230027	128.25	Fantasy	Las Vegas
892230030	152.55	Fantasy	Las Vegas
892230031	176.85	Fantasy	Las Vegas
892249054	152.55	Fantasy	Las Vegas
892342865	176.85	Fantasy	Las Vegas
892360244	176.85	Fantasy	Las Vegas
901075247	128.25	Fantasy	Las Vegas
905895086	144.45	Fantasy	Las Vegas
908305901	176.85	Fantasy	Las Vegas

```
15 rows in set (0.00 sec)
```

## Part 2

**\*\*Events and city with Ticket\_price starting from \$1000\*\***

### Default Configuration:

[illegible]

Configuration 1: create index idx1 on Tickets(ticket\_price);

[illegible]

Configuration 2: create index idx1 on Locations(city);

[illegible]

Configuration 3: create index idx1 on Tickets(ticket\_price); create index idx1 on Locations(city);

```
mysql> explain analyze select event_title, city from Events natural join Locations where event_title in (select event_title from Tickets where ticket_price > 1000);
```

```
-----+-----  
| EXPLAIN  
  
-----+-----  
  
-> Nested loop inner join (cost=103742.12 rows=1035963) (actual time=9.733..12.185 rows=48 loops=1)  
   -> Nested loop inner join (cost=192.22 rows=436) (actual time=0.113..2.269 rows=420 loops=1)  
        -> Covering index scan on locations using idx_location_name (location_name=locations.location_name) (cost=0.25 rows=1) (actual time=0.004..0.005 rows=1 loops=418)  
             -> Single-row index lookup on <subquery>> using <auto distinct key> (event_title=Events.event_title) (actual time=0.023..0.023 rows=0 loops=420)  
                  -> Materialize with deduplication (cost=1307.06..1307.06 rows=2376) (actual time=5.488..5.488 rows=48 loops=1)  
                       -> Filter: (Tickets.event_title is not null) (cost=1069.46 rows=2376) (actual time=0.036..0.263 rows=2376 loops=1)  
                              -> Index range scan on Tickets using idx1 over (>1000.00 < ticket_price), with index condition: (Tickets.ticket_price > 1000.00) (cost=1069.46 rows=2376) (actual time=0.028..8.042 rows=2376 loops=1)
```

We are using `ticket_price` with a 'where' clause. When an index is created on `ticket_price`, the database can quickly locate the rows that satisfy this condition. Without an index, the database would need to perform a full table scan, checking every row, which is much slower, especially as the number of rows increases. Thus, the cost for table scan operation (inside subquery) reduced from 3771 to 1069 for a total decrease from 518,110 to 100,024. However, creating an index on `Locations(city)` increased the cost from 518,110 to 537,389 seemingly because it is strongly associated with `event_title`, a primary key. As one might assume, combining these keys improves overall cost from the default (from 518,110 to 103,742) but results in a worse cost than with only the `Tickets(ticket_price)` indexing (from 100,024 to 103,742) because the `Locations(city)` indexing increases the cost. For this reason, we chose to stick with only the `Tickets(ticket_price)` indexing configuration (Configuration 1).

**\*\*Events and number of tickets available for all events in Great Britain\*\***

### Default Configuration:

[illegible]

Configuration 1: create index idx\_country on Locations(country);

```
mysql> explain analyze select count(ticket_id), event_title, event_url from Tickets natural join Events natural join Locations where country = 'GB' group by event_title, location_name;
+-----+-----+-----+-----+-----+-----+
| | | | | | |
+-----+-----+-----+-----+-----+-----+
| EXPLAIN |
+-----+-----+-----+-----+-----+-----+
| | | | | | |
+-----+-----+-----+-----+-----+-----+
| | | | | | |
+-----+-----+-----+-----+-----+-----+
| -> Table scan on <temporary> (actual time=3.429..3.431 rows=7 loops=1) |
| -> Aggregate using temporary table (actual time=3.424..3.424 rows=7 loops=1) |
| -> Nested loop inner join (cost=126.30 rows=457) (actual time=0.082..1.650 rows=1335 loops=1) |
| -> Nested loop inner join (cost=3.25 rows=6) (actual time=0.041..0.121 rows=7 loops=1) |
| -> Covering index lookup on Locations using idx_country (country='GB') (cost=1.06 rows=6) (actual time=0.014..0.021 rows=6 loops=1) |
| -> Index lookup on Events using idx1 (location_name=Locations.location_name) (cost=0.28 rows=1) (actual time=0.015..0.016 rows=1 loops=6) |
| -> Covering index lookup on Tickets using event_title (event_title=Events.event_title) (cost=13.54 rows=73) (actual time=0.044..0.201 rows=219 loops=7) |
+-----+-----+-----+-----+-----+-----+
| | | | | | |
+-----+-----+-----+-----+-----+-----+
1 row in set (0.01 sec)
```

Configuration 2: create index idx\_url on Events(event\_url);

```
mysql> explain analyze
-> select count(ticket_id), event_title, event_url from Tickets natural join Events natural join Locations where country = 'GB' group by event_title, location_name;
+-----+
|      |
+-----+
| EXPLAIN |
+-----+
|         |
+-----+
|         |
+-----+
|         |
+-----+
|         |
+-----+
|-> Table scan on <temporary> (actual time=3.429..3.430 rows=7 loops=1)
   -> Aggregate using temporary table (actual time=3.427..3.427 rows=7 loops=1)
       -> Nested loop inner join (cost=915.58 rows=3185) (actual time=0.147..1.739 rows=1535 loops=1)
           -> Nested loop inner join (cost=58.30 rows=44) (actual time=0.127..0.329 rows=7 loops=1)
               -> Filter: (Locations.country = 'en') (cost=43.05 rows=42) (actual time=0.038..0.245 rows=6 loops=1)
                   -> Table scan on Locations (cost=43.05 rows=418) (actual time=0.074..0.199 rows=418 loops=1)
                       -> Index lookup on Events using idx1 (location_name=Locations.location_name) (cost=0.26 rows=1) (actual time=0.013..0.014 rows=1 loops=6)
                           -> Covering index lookup on Tickets using event_title (event_title=Events.event_title) (cost=12.54 rows=73) (actual time=0.039..0.161 rows=219 loops=7)

```

Configuration 3: create index idx\_url on Events(event\_url); create index idx\_country on Locations(country);

```
mysql> explain analyze select count(ticket_id), event_title, event_url from Tickets natural join Events natural join Locations where country='GB' group by event_title, location_name;
```

```
-----+-----
```

```
| EXPLAIN
```

```
-----+-----
```

```
|
```

```
-----+-----
```

```
| -> Table scan on temporary< (actual time=3.245..3.246 rows=7 loops=1)
```

```
|   -> Aggregate using temporary table (actual time=3.242..3.242 rows=7 loops=1)
```

```
|     -> Nested loop inner join (cost=126.30 rows=457) (actual time=0.059..1.551 rows=1535 loops=1)
```

```
|       -> Nested loop inner join (cost=3.25 rows=6) (actual time=0.040..0.097 rows=7 loops=1)
```

```
|         -> Covering index lookup on Locations using idx_country (country='GB') (cost=1.06 rows=6) (actual time=0.014..0.017 rows=6 loops=1)
```

```
|           -> Index lookup on Events using idx_l1 (location_name=Locations.location_name) (cost=0.28 rows=1) (actual time=0.012..0.013 rows=1 loops=6)
```

```
|             -> Covering index lookup on Tickets using event_title (event_title=Events.event_title) (cost=13.54 rows=73) (actual time=0.042..0.189 rows=219 loops=7)
```

```
|
```

```
-----+-----
```

```
1 row in set (0.01 sec)
```

One of the query includes a WHERE clause that filters results based on country = 'GB'. If there is an index on the country column, the database can quickly locate rows that match this condition.

**\*\*Events, location, city, number of users who have wishlisted events for the first quarter of 2025\*\***

```
mysql> explain analyze
-> select event_title, location_name, city, count(username) from WishList natural join Events natural join Locations where wishlist_date between '2025-01-01' and '2025-03-01' group by event_title, location_name, city;
+-----+
| EXPLAIN
+-----+
|
+-----+
|
+-----+
| -> Table scan on <temporary> (actual time=3.421..3.472 rows=190 loops=1)
|   -> Aggregate using temporary table (actual time=3.418..3.418 rows=190 loops=1)
|     -> Nested loop inner join (cost=257.24 rows=159) (actual time=0.133..2.901 rows=246 loops=1)
|       -> Nested loop inner join (cost=201.52 rows=159) (actual time=0.126..2.095 rows=246 loops=1)
|         -> Filter: (WishList.wishlist_date between '2025-01-01' and '2025-03-01') (cost=145.80 rows=159) (actual time=0.096..1.226 rows=246 loops=1)
|           -> Table scan on WishList (cost=145.80 rows=1439) (actual time=0.066..0.691 rows=1439 loops=1)
|             -> Filter: ("Events".location_name is not null) (cost=0.25 rows=1) (actual time=0.003..0.003 rows=1 loops=246)
|               -> Single-row index lookup on Events using PRIMARY (event_title=WishList.event_title) (cost=0.25 rows=1) (actual time=0.003..0.003 rows=1 loops=246)
|             -> Single-row index lookup on Locations using PRIMARY (location_name='Events'.location_name) (cost=0.25 rows=1) (actual time=0.003..0.003 rows=1 loops=246)
```

```
mysql> explain analyze select event_title, location_name, city, count(username) from WishList natural join Events natural join Locations where wishlist_date between '2025-01-01' and '2025-03-01' group by event_title, location_name, city;
```

```
+-----+
|
```

```
| EXPLAIN
```

```
+-----+
```

```
|
```

```
+-----+
```

```
| -> Table scan on <temporary> (actual time=2.141..2.171 rows=190 loops=1)
```

```
-> Aggregate using temporary table (actual time=2.139..2.139 rows=190 loops=1)
```

```
-> Nested loop inner join (cost=256.97 rows=246) (actual time=0.054..1.759 rows=246 loops=1)
```

```
-> Nested loop inner join (cost=170.87 rows=246) (actual time=0.045..1.017 rows=246 loops=1)
```

```
-> Filter: (WishList.wishlist_date between '2025-01-01' and '2025-03-01') (cost=84.77 rows=246) (actual time=0.021..0.238 rows=246 loops=1)
```

```
-> Covering index range scan on WishList using idx1 over ('2025-01-01 00:00:00' <= wishlist_date <= '2025-03-01 00:00:00') (cost=84.77 rows=246) (actual time=0.018..0.140 rows=246 loops=1)
```

```
-> Filter: ('Events'.location_name is not null) (cost=0.25 rows=1) (actual time=0.003..0.003 rows=1 loops=246)
```

```
-> Single-row index lookup on Events using PRIMARY (event_title=WishList.event_title) (cost=0.25 rows=1) (actual time=0.003..0.003 rows=1 loops=246)
```

```
-> Single-row index lookup on Locations using PRIMARY (location_name='Events'.location_name) (cost=0.25 rows=1) (actual time=0.003..0.003 rows=1 loops=246)
```

Configuration 2: create index idx2 on WishList(event\_title, wishlisted\_date);

```
mysql> explain analyze  
-> select event_title, location_name, city, count(username) from WishList natural join Events natural join Locations where wishlist_date between '2025-01-01' and '2025-03-01' group by event_title, location_name,  
city;  
+-----+  
| |  
+-----+  
| EXPLAIN |  
+-----+  
| |  
+-----+  
| |  
+-----+  
  
+-----+  
|-> Group aggregate: count(WishList.username) (cost=273.17 rows=159) (actual time=0.244..2.769 rows=190 loops=1)  
|-> Nested loop inner join (cost=257.24 rows=159) (actual time=0.213..2.597 rows=246 loops=1)  
|-> Nested loop inner join (cost=201.52 rows=159) (actual time=0.201..1.900 rows=246 loops=1)  
|-> Filter: (WishList.wishlist_date between '2025-01-01' and '2025-03-01') (cost=145.80 rows=159) (actual time=0.093..0.974 rows=246 loops=1)  
|-> Covering index scan on WishList using idx2 (cost=145.80 rows=1433) (actual time=0.080..0.517 rows=1433 loops=1)  
|-> Filter: ('Events'.location_name is not null) (cost=0.25 rows=1) (actual time=0.003..0.003 rows=1 loops=246)  
|-> Single-row index lookup on Events using PRIMARY (event_title=WishList.event_title) (cost=0.25 rows=1) (actual time=0.003..0.003 rows=1 loops=246)  
|-> Single-row index lookup on Locations using PRIMARY (location_name='Events'.location_name) (cost=0.25 rows=1) (actual time=0.003..0.003 rows=1 loops=246)  
|  
+-----+
```

Configuration 3: create index idx3 on WishList(username, wishlisted\_date);

```
mysql> explain analyze select event_title, location_name, city, count(username) from WishList natural join Events natural join Locations where wishlist_date between '2025-01-01' and '2025-03-01' group by event_title,  
location_name, city;
```

```
+-----+  
|               |  
+-----+  
  
+-----+  
| EXPLAIN        |  
+-----+  
  
+-----+  
|               |  
+-----+  
  
+-----+  
| Group aggregate: count(WishList.username)      (cost=272.67 rows=159) (actual time=0.614..7.510 rows=190 loops=1)  
-> Nested loop inner join (cost=256.74 rows=159) (actual time=0.576..7.289 rows=246 loops=1)  
   -> Nested loop inner join (cost=201.02 rows=159) (actual time=0.560..6.425 rows=246 loops=1)  
       -> Filter: (Wishlist.wishlist_date between '2025-01-01' and '2025-03-01') (cost=145.30 rows=159) (actual time=0.525..5.388 rows=246 loops=1)  
           -> Index scan on Wishlist using event title (cost=145.30 rows=1433) (actual time=0.514..4.811 rows=1433 loops=1)  
               -> Filter: ('Events'.location_name is not null) (cost=0.25 rows=1) (actual time=0.004..0.004 rows=1 loops=246)  
                   -> Single-row index lookup on Events using PRIMARY (event_title=Wishlist.event_title) (cost=0.25 rows=1) (actual time=0.004..0.004 rows=1 loops=246)  
                       -> Single-row index lookup on Locations using PRIMARY (location_name='Events'.location_name) (cost=0.25 rows=1) (actual time=0.003..0.003 rows=1 loops=246)
```

```
|  
+-----+  
|               |  
+-----+
```

```
1 row in set (0.01 sec)
```

We aim to extract the wishlisted events between a specific time frame(quarterly). We extract all the events wishlisted between start date and end date. We also observed that the performance improved given that:

1. There are at least several different values in the date column.
2. Date range used would select less than the most number of rows.

Thus, indexing on wishedlisted\_date(with appropriate time frame selection) reduced the cost from 145.80 to 84.77 on Wishlist table scan operation (Config 1). The overall cost was reduced from



**\*\*Get all Tickets Details, Event details of all events for the top 5 major cities (top 5 cities by events)\*\***

[illegible][illegible]

Configuration 2: create index idx\_price on Tickets(total\_price);

```
mysql> explain analyze
-> Select ticket_id, total price, event_title, city from Tickets natural join Events natural join Locations where city in (Select city from (select city, count(event_title) as Locations_natural_join_Events_group
up by city order by 2 desc limit 5) z);
```

```
+-----+
|EXPLAIN|
+-----+
```

```
+-----+
|Nested loop inner join (cost=813.10 rows=0) (actual time=2.183..20.386 rows=6477 loops=1)|
|-- Nested loop inner join (cost=5.84 rows=0) (actual time=1.894..2.737 rows=78 loops=1)|
|   -- Filter: (Locations.city = <subquery>2).city) (cost=6.28 rows=0) (actual time=1.877..2.258 rows=73 loops=1)|
|     -- Inner hash join (<hash>(Locations.city)<-><hash>(<subquery>2.city)) (cost=5.28 rows=0) (actual time=1.874..2.223 rows=73 loops=1)|
|       -- Table scan on Locations (cost=5.46 rows=410) (actual time=0.746..0.269 rows=418 loops=1)|
|         -- Hash |
|           -- Table scan on <subquery>2 (cost=5.00..5.00 rows=0) (actual time=1.793..1.793 rows=5 loops=1)|
|             -- Materialize with deduplication (cost=2.50..2.50 rows=0) (actual time=1.792..1.792 rows=5 loops=1)|
|               -- Filter: (<z.city is not null) (cost=2.50..2.50 rows=0) (actual time=1.781..1.782 rows=5 loops=1)|
|                 -- Table scan on z (cost=2.50..2.50 rows=0) (actual time=1.779..1.780 rows=5 loops=1)|
|                   -- Materialize (cost=0.00..0.00 rows=0) (actual time=1.779..1.779 rows=5 loops=1)|
|                     -- Limit: 5 row(s) (actual time=1.767..1.768 rows=5 loops=1)|
|                       -- Sort: count(event_title) DESC, limit input to 5 row(s) per chunk (actual time=1.767..1.767 rows=5 loops=1)|
|                         -- Table scan on <temporary> (actual time=1.696..1.724 rows=222 loops=1)|
|                           -- Aggregate using temporary table (actual time=1.693..1.693 rows=222 loops=1)|
|                             -- Nested loop inner join (cost=190.75 rows=420) (actual time=0.117..1.384 rows=420 loops=1)|
|                               -- Filter: ('Events.location name' is not null) (cost=43.75 rows=420) (actual time=0.088..0.259 rows=420 loops=1)|
|                                 -- Covering index scan on Events using index (cost=43.75 rows=420) (actual time=0.086..0.224 rows=420 loops=1)|
|                                   -- Single-row index lookup on Locations using PRIMARY (location name='Events'.location name) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=420)|
|                                     -- Index lookup on Events using event_title (event_title='Events'.event_title) (cost=0.26 rows=1) (actual time=0.005..0.006 rows=1 loops=73)|
|                                       -- Index lookup on Tickets using event_title (event_title='Events'.event_title) (cost=18.45 rows=3) (actual time=0.886..0.221 rows=3 loops=18)
```

Configuration 3: create index idx1 on Locations(city); create index idx\_price on Tickets(total\_price);

[illegible]

We aim to extract the top 5 cities by events and display tickets associated with these cities. Thus, we add an index to the city column, which is used with a 'where' clause. We observe that the number of operations for the query is reduced after indexing. Before indexing, the query scanned the entire table Location and performed an inner hash join with the subquery. After indexing, index lookup was performed initially which resulted in reduction of cost of overall query from 813.10 to 113.09 (Config 1). With only Tickets(total\_price) as index, the overall cost remained constant (Config 2). When both the previous indexes are considered, the cost is reduced to 113.09 (Config 3 same as Config 1). Thus we went ahead with Config 1, since indexing on total\_price did not affect the query performance.

### **Final Index Design:**

- create index idx1 on Locations(city);
- create index idx1 on WishList(wishlisted\_date);
- create index idx\_country on Locations(country);
- create index idx1 on Tickets(ticket\_price);